

# Assignment\_3\_Report

*Jakub Domasik*

*7 12 2019*

## Q1

(a)

Here is the provided gradient descent function.

```
gradient.descent <- function(f, gradf, x0, iterations=1000, eta=0.2) {  
  x<-x0  
  for (i in 1:iterations) {  
    if(i<50) {  
      cat(i,"/",iterations,": ",x," ",f(x),"\n")  
    }  
    x<-x-eta*gradf(x)  
  }  
  x  
}
```

```
gradient.ascent <- function(f, df, x0, iterations=1000, eta=0.2) {  
  gradient.descent(f, df, x0, iterations, -eta)  
}
```

To test the function, the following code is used.

```
f <-function(x) { (1+x^2)^(-1) }  
gradf<-function(x) { -2*x*(1+x^2)^(-2) }  
gradient.ascent(f,gradf,3,40,0.5)
```

```
## 1 / 40 : 3 0.1  
## 2 / 40 : 2.97 0.1018237  
## 3 / 40 : 2.939207 0.1037459  
## 4 / 40 : 2.907572 0.1057756  
## 5 / 40 : 2.87504 0.1079231  
## 6 / 40 : 2.841554 0.1101998  
## 7 / 40 : 2.807046 0.1126189  
## 8 / 40 : 2.771444 0.1151954  
## 9 / 40 : 2.734667 0.1179467  
## 10 / 40 : 2.696624 0.120893  
## 11 / 40 : 2.657212 0.1240575  
## 12 / 40 : 2.616317 0.1274679  
## 13 / 40 : 2.573807 0.1311564  
## 14 / 40 : 2.529532 0.1351619  
## 15 / 40 : 2.483321 0.1395307  
## 16 / 40 : 2.434974 0.144319  
## 17 / 40 : 2.384258 0.1495956  
## 18 / 40 : 2.330901 0.155446
```

```

## 19 / 40 : 2.274579 0.1619772
## 20 / 40 : 2.214901 0.1693254
## 21 / 40 : 2.151398 0.1776669
## 22 / 40 : 2.083488 0.1872336
## 23 / 40 : 2.010448 0.1983379
## 24 / 40 : 1.931361 0.2114095
## 25 / 40 : 1.845041 0.2270572
## 26 / 40 : 1.74992 0.2461708
## 27 / 40 : 1.643875 0.2701006
## 28 / 40 : 1.523947 0.300986
## 29 / 40 : 1.385889 0.3423852
## 30 / 40 : 1.223424 0.400518
## 31 / 40 : 1.027169 0.4866
## 32 / 40 : 0.7839563 0.6193532
## 33 / 40 : 0.4832319 0.8106927
## 34 / 40 : 0.165641 0.9732957
## 35 / 40 : 0.008728518 0.9999238
## 36 / 40 : 1.329848e-06 1
## 37 / 40 : 4.703997e-18 1
## 38 / 40 : 0 1
## 39 / 40 : 0 1
## 40 / 40 : 0 1

```

```
## [1] 0
```

(b) (i)

We shall prove that function  $f(x_1, x_2) = (x_1 - 1)^2 + 100(x_1^2 - x_2)^2$  has one unique stationary point which is a minimum.

1. Find stationary points.

(a) Calculate derivatives of  $f(x_1, x_2)$ .

$$\frac{df}{dx_1} = 2(x_1 - 1) + 4 * 100x_1(x_1 - x_2) = 2x_1 - 2 + 400x_1^3 - 400x_1x_2$$

$$\frac{df}{dx_2} = 200x_1^2 + 200x_2$$

(b) Set the derivatives equal to 0.

$$(1) 2x_1 - 2 + 400x_1^3 - 400x_1x_2 = 0$$

$$(2) 200x_1^2 + 200x_2 = 0$$

(c) Solve the system of two equations.

$$(2) \implies x_2 = x_1^2$$

$$(1) \implies 2x_1 - 2 + 400x_1^3 - 400x_1^3 \implies x_1 = 1 \implies x_2 = 1$$

Therefore, we found a unique stationary point at  $(1, 1)$

2. Check what kind of stationary point is  $(1, 1)$ .

(a) Calculate 2<sup>nd</sup> derivatives.

$$\frac{d^2 f}{dx_1^2} = 2 + 1200x_1 + 400x_2$$

$$\frac{d^2 f}{dx_2^2} = 200$$

$$\frac{d^2 f}{dx_1 dx_2} = -400x_1$$

$$f(1, 1) = 1602 > 0$$

$$\det(D^2 f) = f_{x_1 x_1} * f_{x_2 x_2} - f_{x_1 x_2}^2 = 160400$$

$D^2 f(1, 1)$  is positive definite so  $(1, 1)$  is a minimum. Therefore, we proved that the provided function  $f$  has a unique minimum.

(b)(ii)

For the function  $f$  as defined above, gradient of  $f$  is returned using *gradient\_f* function

```
f <- function(x) (x[1]-1)^2 + 100*(x[1]^2-x[2])^2
gradient_f <- function(x) {
  return (c( 2 * x[1] - 2 + 400 * x[1]^3 - 400 * x[2] * x[1],
            -200 * x[1]^2 + 200 * x[2]))
}
```

(b) (iii)

Gradient descent function is used to find the minimum of the function  $f$ . I am asked to start the gradient algorithm at  $x_0 = (3, 4)$ . We set number of iterations to 30 000 and  $\alpha = 0.0007$ . The converges to 0 as  $(x_1, x_2)$  goes to  $(1, 1)$ .

```
gradient.descent(f, gradient_f, c(3, 4), 30000, 0.0007)
```

```
## 1 / 30000 : 3 4 2504
## 2 / 30000 : -1.2028 4.7 1063.23
## 3 / 30000 : -2.295366 4.244542 115.7505
## 4 / 30000 : -1.63252 4.387925 303.7352
## 5 / 30000 : -2.416337 4.146732 297.9424
## 6 / 30000 : -1.266821 4.383606 777.2945
## 7 / 30000 : -2.249305 3.994578 123.9369
## 8 / 30000 : -1.574142 4.14365 284.0905
## 9 / 30000 : -2.304723 3.910448 207.286
## 10 / 30000 : -1.395805 4.10663 471.5913
## 11 / 30000 : -2.235992 3.80446 153.3222
## 12 / 30000 : -1.483173 3.971788 320.1597
## 13 / 30000 : -2.215582 3.72371 150.7845
```

```
## 14 / 30000 : -1.475892 3.889623 299.007
## 15 / 30000 : -2.179647 3.650032 131.293
## 16 / 30000 : -1.503357 3.804148 244.6803
## 17 / 30000 : -2.149811 3.587979 116.777
## 18 / 30000 : -1.523163 3.732699 205.9311
## 19 / 30000 : -2.122115 3.534924 103.537
## 20 / 30000 : -1.542299 3.670507 173.3432
## 21 / 30000 : -2.096605 3.489652 91.6904
## 22 / 30000 : -1.560345 3.616506 146.2276
## 23 / 30000 : -2.073098 3.45105 81.13146
## 24 / 30000 : -1.577323 3.569586 123.637
## 25 / 30000 : -2.051421 3.418156 71.74812
## 26 / 30000 : -1.593276 3.52878 104.7846
## 27 / 30000 : -2.031414 3.390145 63.43223
## 28 / 30000 : -1.608253 3.493255 89.02718
## 29 / 30000 : -2.012933 3.366306 56.08149
## 30 / 30000 : -1.6223 3.462289 75.83805
## 31 / 30000 : -1.995847 3.346029 49.60019
## 32 / 30000 : -1.635463 3.435262 64.78533
## 33 / 30000 : -1.980039 3.328788 43.89947
## 34 / 30000 : -1.647785 3.411636 55.51369
## 35 / 30000 : -1.965402 3.314134 38.89742
## 36 / 30000 : -1.65931 3.390948 47.72999
## 37 / 30000 : -1.951838 3.301679 34.51896
## 38 / 30000 : -1.67008 3.372798 41.19174
## 39 / 30000 : -1.939261 3.291089 30.69561
## 40 / 30000 : -1.680133 3.356839 35.6978
## 41 / 30000 : -1.927591 3.282081 27.36518
## 42 / 30000 : -1.689508 3.342774 31.08089
## 43 / 30000 : -1.916757 3.274407 24.47142
## 44 / 30000 : -1.698238 3.330344 27.20154
## 45 / 30000 : -1.906694 3.267857 21.96359
## 46 / 30000 : -1.70636 3.319325 23.9432
## 47 / 30000 : -1.897344 3.262252 19.79607
## 48 / 30000 : -1.713903 3.309524 21.20821
## 49 / 30000 : -1.888653 3.257436 17.92792
```

```
## [1] 0.9998311 0.9996615
```

(c)

The function *gradient.momentum* performs the algorithm: gradient descent with momentum. We use it to find the minimum of the function *f* defined in part(b).

```
gradient.momentum <- function(f, gradf, x0, iterations = 1000, eta = 0.2, alpha) {
  x1 <- x0
  x2 <- x1 - eta * gradf(x1)
  cat(1, " : ", x1, ", ", f(x1), "\n")
  cat(2, " : ", x2, ", ", f(x2), "\n")
  for( i in 3:iterations) {
    x <- x2 - eta * gradf(x2) + alpha * (x2 - x1)
    x1 <- x2
    x2 <- x
  }
}
```

```

    if(i < 50) {
      cat(i, " : ", x, ", ", f(x), "\n")
    }
  }
  x

}

gradient.momentum(f, gradient_f, c(3,4), 30000, 0.0005, 0.03)

```

```

## 1 : 3 4 , 2504
## 2 : -0.002 4.5 , 2026
## 3 : -0.092858 4.065 , 1646.614
## 4 : -0.1698243 3.646313 , 1309.979
## 5 : -0.2938304 3.272005 , 1016.522
## 6 : -0.483466 2.942209 , 735.7813
## 7 : -0.7495622 2.661468 , 443.9032
## 8 : -1.070555 2.443083 , 172.5067
## 9 : -1.355815 2.306832 , 27.50815
## 10 : -1.489084 2.255885 , 6.343879
## 11 : -1.502063 2.250505 , 6.263555
## 12 : -1.498242 2.250913 , 6.245036
## 13 : -1.497482 2.250306 , 6.243585
## 14 : -1.497314 2.249503 , 6.242283
## 15 : -1.497074 2.248723 , 6.240992
## 16 : -1.496813 2.24795 , 6.239702
## 17 : -1.496554 2.247177 , 6.238412
## 18 : -1.496295 2.246404 , 6.237121
## 19 : -1.496037 2.24563 , 6.235831
## 20 : -1.495778 2.244856 , 6.23454
## 21 : -1.49552 2.244083 , 6.233249
## 22 : -1.495261 2.243309 , 6.231958
## 23 : -1.495002 2.242536 , 6.230667
## 24 : -1.494743 2.241762 , 6.229376
## 25 : -1.494484 2.240988 , 6.228084
## 26 : -1.494225 2.240214 , 6.226793
## 27 : -1.493966 2.239441 , 6.225501
## 28 : -1.493707 2.238667 , 6.224209
## 29 : -1.493448 2.237893 , 6.222917
## 30 : -1.493189 2.237119 , 6.221625
## 31 : -1.49293 2.236345 , 6.220333
## 32 : -1.49267 2.235572 , 6.219041
## 33 : -1.492411 2.234798 , 6.217748
## 34 : -1.492152 2.234024 , 6.216455
## 35 : -1.491892 2.23325 , 6.215163
## 36 : -1.491633 2.232476 , 6.21387
## 37 : -1.491373 2.231702 , 6.212577
## 38 : -1.491113 2.230928 , 6.211284
## 39 : -1.490854 2.230154 , 6.20999
## 40 : -1.490594 2.229379 , 6.208697
## 41 : -1.490334 2.228605 , 6.207403
## 42 : -1.490074 2.227831 , 6.20611
## 43 : -1.489814 2.227057 , 6.204816

```

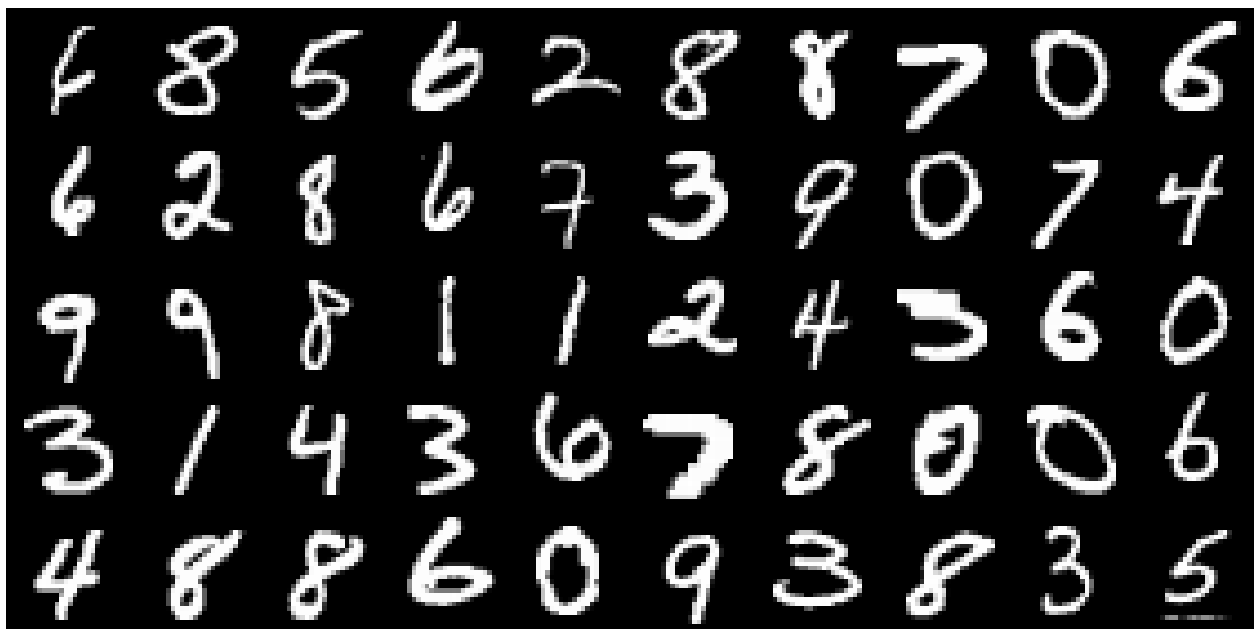
```
## 44 : -1.489554 2.226283 , 6.203522
## 45 : -1.489294 2.225509 , 6.202228
## 46 : -1.489034 2.224734 , 6.200933
## 47 : -1.488774 2.22396 , 6.199639
## 48 : -1.488514 2.223186 , 6.198345
## 49 : -1.488254 2.222411 , 6.19705
```

```
## [1] 0.9988378 0.9976723
```

## Q2

(a)

```
load("mnist.tiny.RData")
train.X=train.X/255
test.X=test.X/255
library(grid)
grid.raster(array(aperm(array(train.X[1:50,],c(5,10,28,28)),c(4,1,3,2)),c(140,280)),
             interpolate=FALSE)
```



```
library(e1071)
```

At first, we run SVM with the linear kernel.

```
svm(train.X,train.labels,type="C-classification",kernel="linear",cross=3)$tot.accuracy
```

```
## [1] 85.8
```

Now, we use the polynomial kernel and try different degrees to find the optimal value.

```
svm(train.X,train.labels,type="C-classification",kernel="poly",degree=2,coef=1,cross=3)$tot.accuracy
```

```
## [1] 82
```

```
svm(train.X,train.labels,type="C-classification",kernel="poly",degree=5,coef=1,cross=3)$tot.accuracy
```

```
## [1] 86.6
```

```
svm(train.X,train.labels,type="C-classification",kernel="poly",degree=10,coef=1,cross=3)$tot.accuracy
```

```
## [1] 88.1
```

```
svm(train.X,train.labels,type="C-classification",kernel="poly",degree=13,coef=1,cross=3)$tot.accuracy
```

```
## [1] 88.4
```

```
svm(train.X,train.labels,type="C-classification",kernel="poly",degree=14,coef=1,cross=3)$tot.accuracy
```

```
## [1] 89.6
```

```
svm(train.X,train.labels,type="C-classification",kernel="poly",degree=15,coef=1,cross=3)$tot.accuracy
```

```
## [1] 88.6
```

```
svm(train.X,train.labels,type="C-classification",kernel="poly",degree=16,coef=1,cross=3)$tot.accuracy
```

```
## [1] 88.4
```

The SVM with linear kernel has the accuracy of 85.8% and performs better than SVM with polynomial kernel of 1st degree. SVM performs the best overall with polynomial kernel of 13th degree.

Now, we try the radial kernel for the SVM and look for the optimal gamma.

```
set.seed(1)
svm(train.X, train.labels, type = "C-classification", kernel = "radial",
     gamma = 1, coef = 1, cross = 3)$tot.accuracy
```

```
## [1] 12.4
```

```
set.seed(1)
svm(train.X, train.labels, type = "C-classification", kernel = "radial",
     gamma = 0.5, coef = 1, cross = 3)$tot.accuracy
```

```
## [1] 13.5
```

```
set.seed(1)
svm(train.X, train.labels, type = "C-classification", kernel = "radial",
     gamma = 0.01, coef = 1, cross = 3)$tot.accuracy
```

```
## [1] 89
```

```
set.seed(1)
svm(train.X, train.labels, type = "C-classification", kernel = "radial",
     gamma = 0.03, coef = 1, cross = 3)$tot.accuracy
```

```
## [1] 90.9
```

```
set.seed(1)
svm(train.X, train.labels, type = "C-classification", kernel = "radial",
     gamma = 0.02, coef = 1, cross = 3)$tot.accuracy
```

```
## [1] 90.7
```

We see that the optimal value of gamma in this case is 0.03. For this value of gamma SVM with radial kernel has higher accuracy than any of SVMs with polynomial or linear kernels.

(b)

```
log.C.range <- log(c(0.001, 0.01, 0.1, 1, 10, 100))
log.gamma.range <- log(c(0.001, 0.01, 0.1, 1, 10, 100))

comb <- matrix(ncol = 2, nrow = length(log.gamma.range) * length(log.C.range))
count <- 0
results <- c()

for(i in 1:length(log.C.range)) {
  for(j in 1:length(log.gamma.range)) {
    count <- count + 1
    comb[count, ] <- c(gamma = exp(log.C.range[i]), exp(log.gamma.range[j]))
  }
}

for ( i in 1:nrow(comb)) {
  accuracy <- svm(train.X, train.labels, type = "C-classification", kernel = "rad",
                  cost = comb[i,1], gamma = comb[i,2], degree = 2, coef = 1, cross = 3)$tot.accuracy
  results <- c(results, accuracy)
```



```
}
row <- which.max(results)
comb[row, ]
```

```
## [1] 10.00 0.01
```

```
max(results)
```

```
## [1] 90.8
```

As calculated above, the optimal values for our model are:  $\text{cost} = 10$ ,  $\text{gamma} = 0.01$ . Now, we train our model on the *tiny* training set and test it on the training set. In the end, we check the accuracy of our model.

```
model <- svm(train.X, train.labels, type = "C-classification", kernel = "rad",
             cost = 10, gamma = 0.01, degree = 2, coef = 1)
test.predictions <- as.vector(predict(model, test.X), mode = "numeric")

accuracy <- sum(diag(table(test.labels, test.predictions))) / sum(table(test.labels, test.predictions))
accuracy
```

```
## [1] 0.913
```

The model got the accuracy of 91.3% on the test set.