

ST340 Lab 5: Bandit problems

2019–20

Bernoulli Bandits

```
library(mvtnorm)
```

```
## Warning: package 'mvtnorm' was built under R version 3.5.2
```

(a) Set Bernoulli success parameters for each arm.

```
ps <- c(0.4,0.6)
```

(b) This is a template for an Epsilon-greedy algorithm, runs for n steps:

```
epsilon.greedy <- function(ps,epsilon,n) {  
  as <- rep(0,n)  
  rs <- rep(0,n)  
  
  ## initial number of plays and number of successes is 0 for each arm  
  ns <- rep(0,2); ss <- rep(0,2)  
  
  ## at first, play each arm once  
  for (i in 1:2) {  
    a <- i  
    r <- runif(1) < ps[a]  
    ns[a] <- ns[a] + 1  
    ss[a] <- ss[a] + r  
    as[i] <- a  
    rs[i] <- r  
  }  
  
  ## now follow the epsilon greedy strategy  
  for (i in 3:n) {  
    # with probability epsilon, pick an arm uniformly at random  
    if (runif(1) < epsilon) {  
      a <- sample(2,1)  
    } else { # otherwise, choose the "best arm so far".  
      a <- which.max(ss/ns)  
    }  
    ## simulate the reward  
    r <- runif(1) < ps[a]  
  
    # update the number of plays, successes  
    ns[a] <- ns[a] + 1  
    ss[a] <- ss[a] + r  
  
    # record the arm played and the reward received  
    as[i] <- a  
    rs[i] <- r  
  }  
}
```

```
    return(list(as=as,rs=rs))
}
```

Run `epsilon.greedy` with the given `ps` and a choice of `epsilon` and see how well it does.

```
eg.out <- epsilon.greedy(ps=ps,epsilon=.1,n=1e4)
```

(c) Implement a `sample_arm` routine, for use in the Thompson sampling code below.

```
sample_arm.bernoulli <- function(ns,ss) {

  ## YOUR CODE HERE

}

thompson.bernoulli <- function(ps,n) {
  as <- rep(0,n)
  rs <- rep(0,n)

  ## number of times each arm has been played
  ## and number of corresponding successes
  ns <- rep(0,2); ss <- rep(0,2)

  for (i in 1:n) {
    a <- sample_arm.bernoulli(ns,ss)
    r <- runif(1) < ps[a]
    ns[a] <- ns[a] + 1
    ss[a] <- ss[a] + r
    as[i] <- a
    rs[i] <- r
  }
  return(list(as=as,rs=rs))
}
```

(d) Run the Thompson scheme and compare its performance to that of `epsilon.greedy`.

```
thompson.bernoulli.out <- thompson.bernoulli(ps=ps,n=1e4)
```